

Working with Multi-Modal Health Data

Matthew Engelhard

Working with categorical variables

One-hot / Cell coding

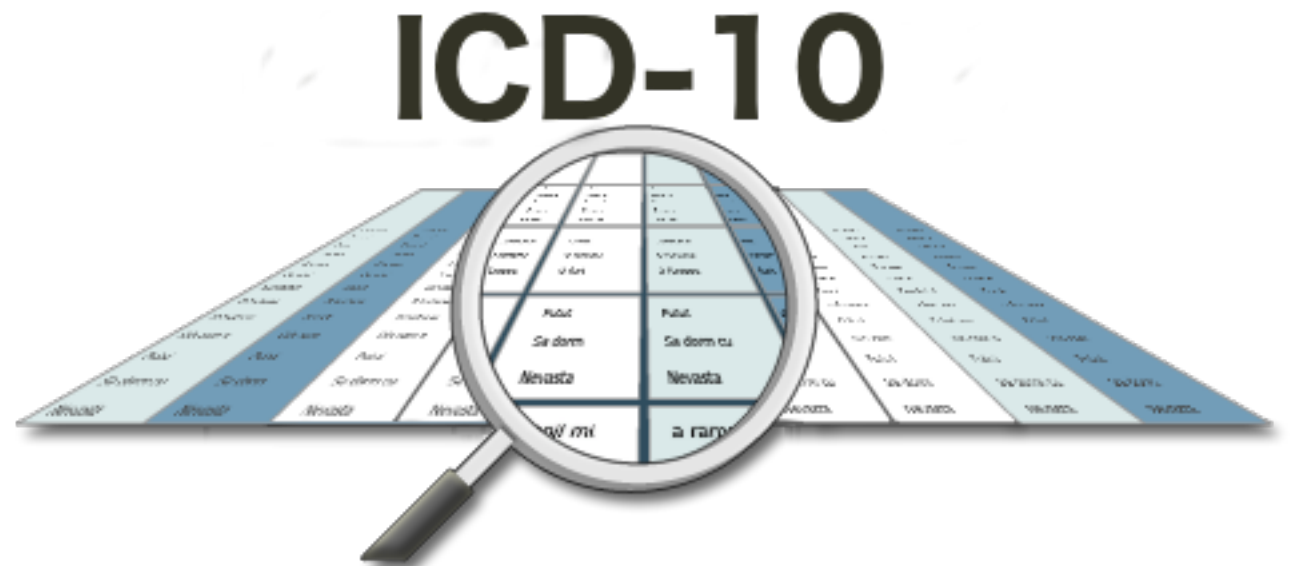
Embeddings

Cell / “One-hot” encoding

- Familiar from linear models: e.g. 0 for male, 1 for female
- Requires # features equal to # categories
- Always categorical: diagnosis codes, procedure codes
- Reduce to categorical: look at which labs are collected, ignore values
- Consider similarity to bag of words models...

Very large number of categories

- Option 1: don't worry about it
- Option 2: group manually or using existing software (e.g. clinical classification software)
- Option 3: use or create embeddings
 - pre-train
 - learn directly



Measurements over time


- Option 1: Count number of occurrences (+ log transform)
 - Option 2: If we have embeddings, we can aggregate (like VSWEM)
 - Option 3: Use a sequential model (covered in the next lecture)
-
- How do we incorporate information about the TIMESTAMP??
 - Positional embedding
 - Temporal embedding

Working with irregularly spaced measurements


Aggregation

Imputation


Measurements in the ICU



Measurements in the ICU




Measurements on the Wards...



This is a major difficulty!


DIHI Sepsis Watch



<- Use GP regression to predict measurements at regular intervals

<- Predict sepsis risk using an RNN

Simplest Method...



aggregate in hour 1 aggregate in hour 2 aggregate in hour 2 aggregate in hour 4...


In the EHR, measurements are highly “sparse”

- Many more missing measurements than non-missing
- Consider diagnosis codes, procedure codes, uncommon labs, etc
- We want to learn from these measurements, but most patients don't have them
- Choose an aggregation window to reduce this, but nevertheless a great deal of imputation is required

Making predictions from multi-modal data

“Data fusion”

Logistic Regression




x_i is a vector of predictor variables

We predict $\text{logit}(p_i)$
(the log-odds)
Directly from x_i

No intermediate or hierarchical
features

Convolutional Neural Network




ω_i is a vector of high-level features extracted by the CNN

We predict $\text{logit}(p_i)$ (the log-odds) based on ω_i

No intermediate or hierarchical features

Any Modality



We predict $\text{logit}(p_i)$
(the log-odds)
based on ω_i


Which may be observed predictors, or
it may be extracted by a deep NN


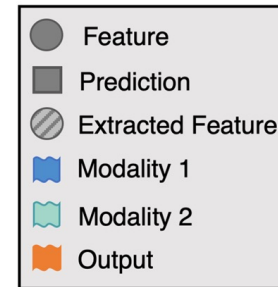




Combining Modalities

1. Train separate models for each modality. This is more important for images, text, and other high-dimensional and/or structured data than for tabular data.
2. Create a single model by concatenating the highest-level features (used to predict $\text{logit}(p_i)$) across all modalities.
3. Consider adding one more fully-connected layer; this allows you to take advantage of interactions between features from different modalities.
4. Train the final layer(s)
5. Consider fine-tuning the entire model

Combining Modalities

1. Train independently
2. Concatenate
3. Train together





Huang, SC., Pareek, A., Seyyedi, S. *et al.* Fusion of medical imaging and electronic health records using deep learning: a systematic review and implementation guidelines. *npj Digit. Med.* **3**, 136 (2020).

Conclusions

- In healthcare, we commonly have the following challenges:
 1. Categorical variables with a large number of categories (i.e. high cardinality)
 2. Sparse measurements
 3. Multiple modalities we need to integrate
- Always start with the simplest method
- Often, aggregating measurements/features performs quite well
- Once our simple baseline is established, we can consider more complex models (e.g. RNNs for sequences)