

# The Model Development and Evaluation Process

Matthew Engelhard

# Today

## Essentials of Sound Model Development and Evaluation

- Why out-of-sample evaluation (*i.e.*, a held-out test set) is critical
- How do we divide up the data?
- How is each portion is used?
- What can go wrong?

## Common Variations in Practice


- Hyperparameter tuning
- Cross-validation

# The Napoleon Dynamite Problem (NDP)

The New York Times Magazine

THE SCREENS ISSUE

## If You Liked This, You're Sure to Love That

 Give this article



By Clive Thompson

Nov. 21, 2008

**THE “NAPOLEON DYNAMITE”** problem is driving Len Bertoni crazy. Bertoni is a 51-year-old “semiretired” computer scientist who lives an hour outside Pittsburgh. In the spring of 2007, his sister-in-law e-mailed him an intriguing bit of news: Netflix, the Web-based DVD-rental company, was holding a contest to try to improve Cinematch, its “recommendation engine.” The prize: \$1 million.



# Claim: I've solved the NDP.

- Randy took headshots of each member of the MMCI class of 2021
- Catherine asked them whether they liked Napoleon Dynamite
- I then studied the data


I can predict whether someone likes Napoleon Dynamite from a picture of their face with 100% accuracy, and I can prove it. Show me a picture of anyone in the class and I will predict whether they like Napoleon Dynamite.

Aren't you impressed?

# You got me. I memorized the dataset.

And I can do it again with this dataset.


I can predict whether people like Napoleon Dynamite from their age and height.



# You got me. I memorized the dataset.

And I can do it again with this dataset.

Question: Can logistic regression memorize this dataset (*i.e.*, make perfect predictions)?

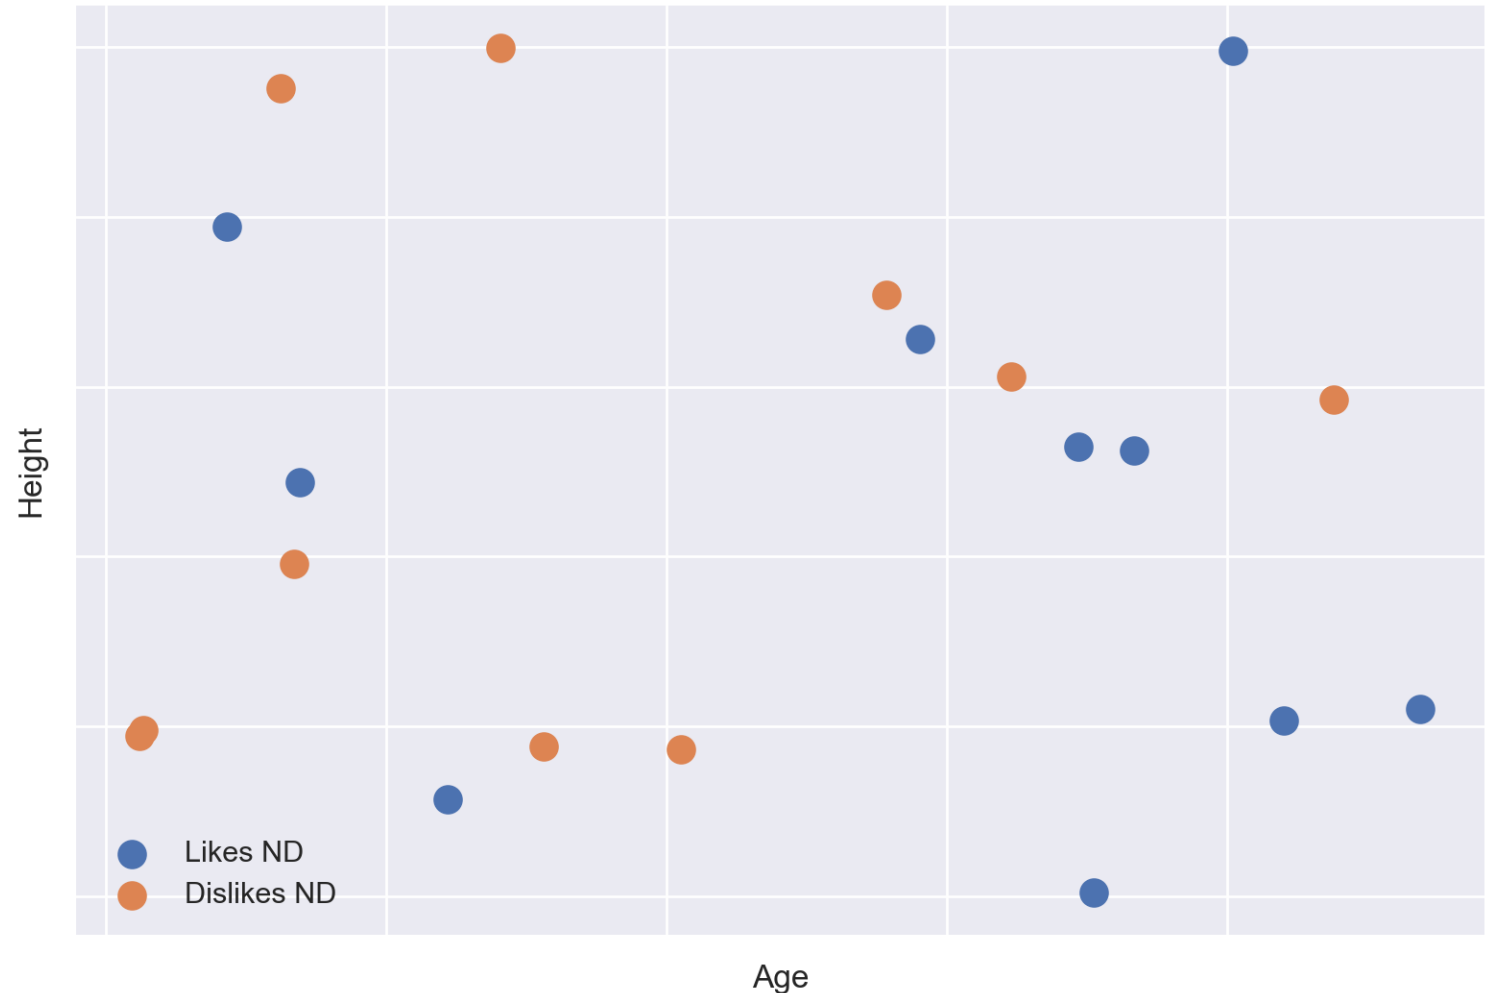


# You got me. I memorized the dataset.

And I can do it again with this dataset.

Question: Can logistic regression memorize this dataset (*i.e.*, make perfect predictions)?


Answer: No. There is no line that perfectly separates the blue and orange points.



# You got me. I memorized the dataset.

And I can do it again with this dataset.

Question: Can a neural network memorize this dataset (*i.e.*, make perfect predictions)?






# You got me. I memorized the dataset.

And I can do it again with this dataset.

Question: Can a neural network memorize this dataset (*i.e.*, make perfect predictions)?

Answer: Yes... like so -->

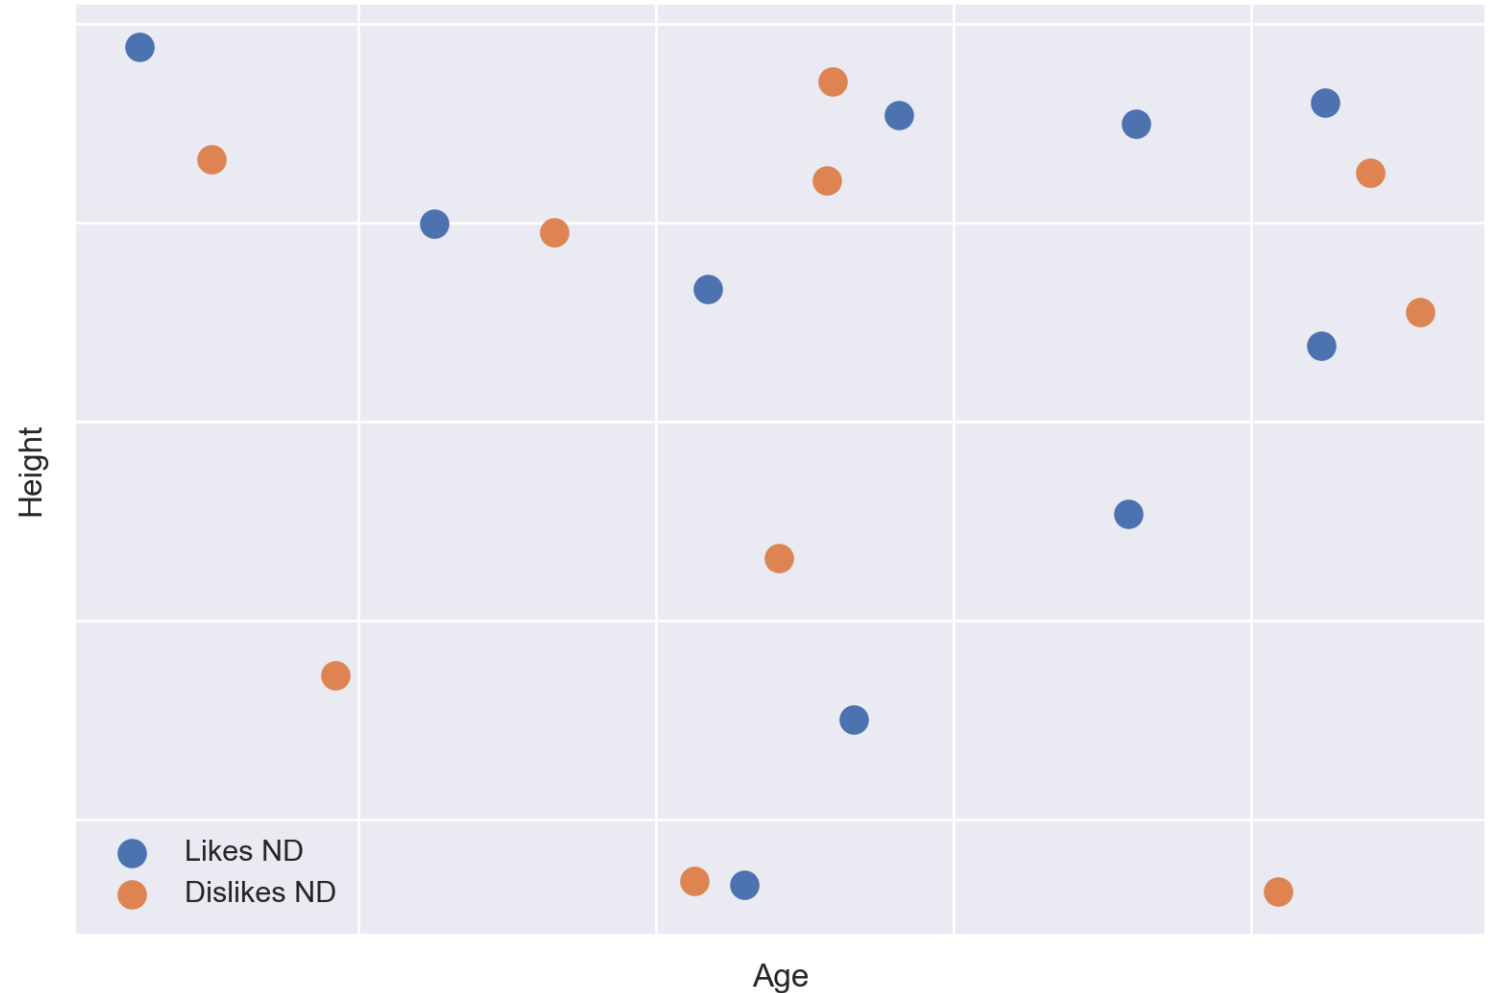


# You got me. I memorized the dataset.

And I can do it again with this dataset.

Question: Can a neural network memorize this dataset (*i.e.*, make perfect predictions)?

Answer: Yes... like so -->




# You got me. I memorized the dataset.

And I can do it again with this dataset.

Question: Can a neural network memorize this dataset (*i.e.*, make perfect predictions)?

Answer: Yes... like so -->




# You got me. I memorized the dataset.

And I can do it again with this dataset.

Question: Can a neural network memorize this dataset (*i.e.*, make perfect predictions)?


Answer: Yes... like so -->

*What is this called?*



You got me. I memorized the dataset.

If this were the true relationship between age, height, and Napoleon Dynamite, then this boundary would still work when we gather data from a new sample of participants.

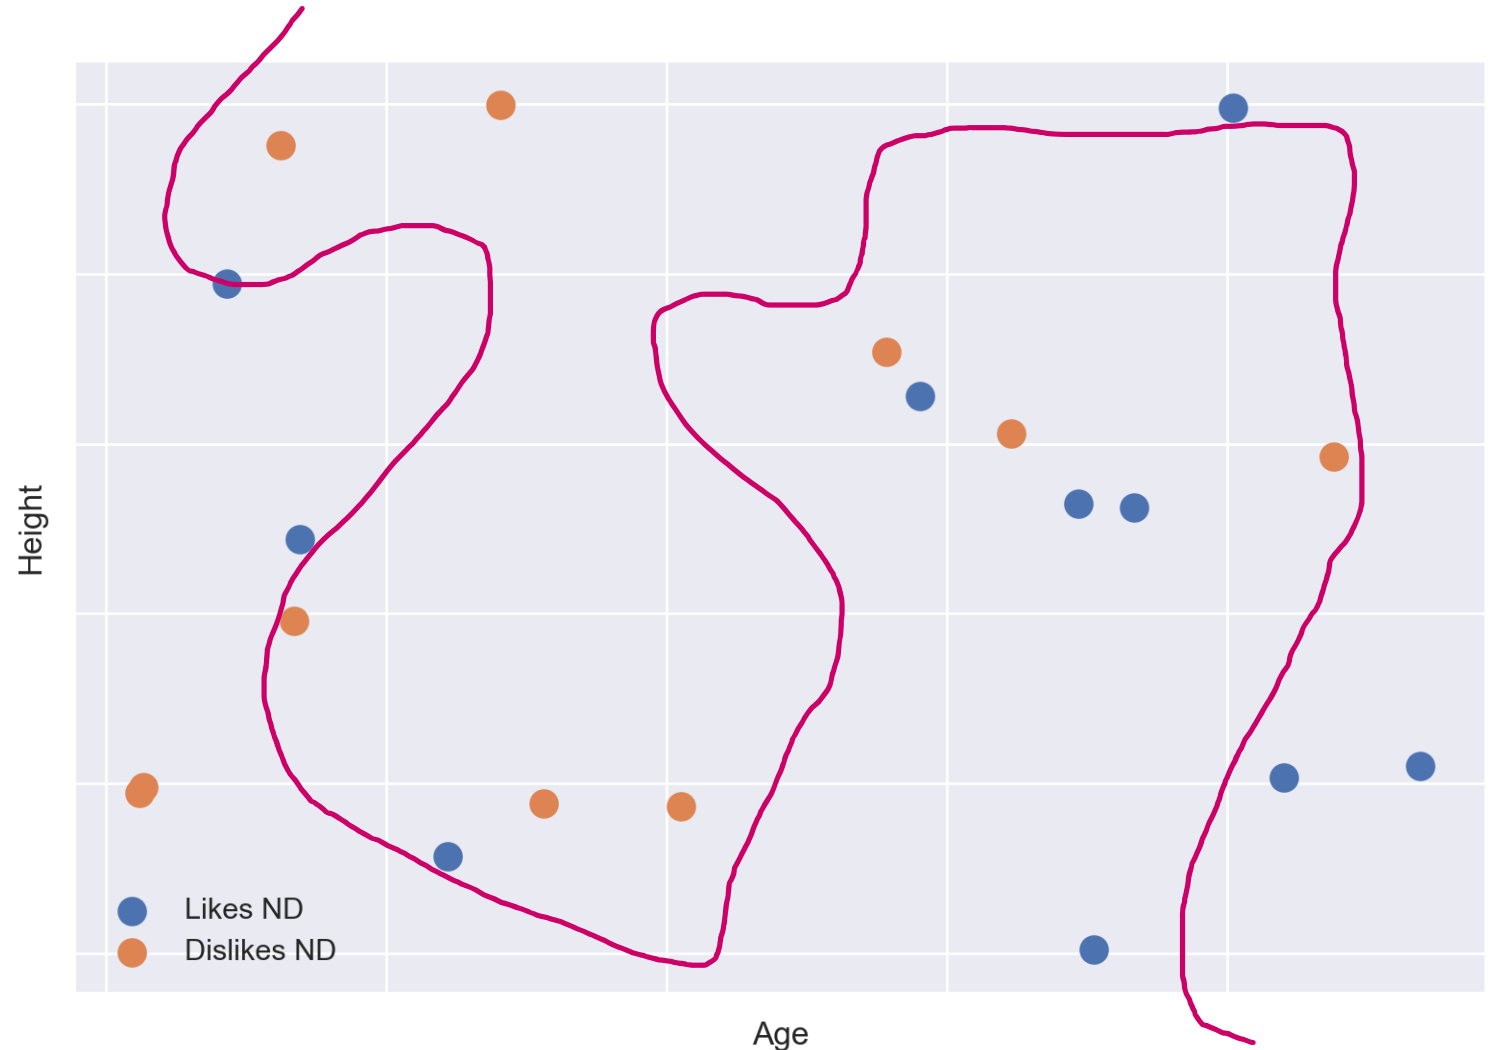


# You got me. I memorized the dataset.

If this were the true relationship between age, height, and Napoleon Dynamite, then this boundary would still work when we gather data from a new sample of participants.

**But it doesn't.**


In this case, there's no relationship – it's just random noise. But the NN can still make perfect predictions on the training data.



# Key Takeaways


- **We must evaluate our models using data they haven't seen before.**
  - The test set should not be used to train
  - The test set should not be used *in any other way* prior to the evaluation
- For complex models – say, an MLP with many layers, or logistic regression with a large number of predictors:
  - Performance on the training set can dramatically overestimate performance on the test set
  - So, evaluating on a held-out test set is critical
- For simple models – say, logistic regression with only a few predictors:
  - Performance on the training set should be close to performance on the test set
  - Nevertheless, evaluating on a held-out test set is good practice

# How do we divide up the data? (easy version)



Development Set:

-> used to build the model



Test Set:

-> used to estimate real-world performance



# Suppose we receive a new dataset

- Read the data (Excel, SQL)
- Assign to the development set or test set: shuffle, then split
- Divide predictors (x) from outcomes (y) and unwanted columns (often identifiers)
- Preprocessing
  - Normalize numeric variables
  - Cell code categorical variables
  - Remember, we don't want to use any information from the test set prior to training. This can be a bit tricky to implement.

MORTALITY PREDICTION WORKSHEET				
COVARIATES				OUTCOME
patient	age	female	temp	mortality
0	30.5	0	105.0	1
1	74.0	1	96.7	0
2	27.4	0	96.1	0
3	0.1	1	98.5	0
4	0.7	1	96.5	0
5	49.9	1	97.1	0
6	72.9	1	100.1	1
7	29.1	1	99.6	0
8	83.5	1	100.6	1
9	82.3	1	95.2	1
10	23.7	0	99.4	1
11	12.9	0	96.6	0
12	53.9	1	100.3	0
13	18.8	0	98.6	0
14	51.8	0	98.5	0
15	3.3	0	94.6	0
16	69.7	0	99.1	0
17	60.4	1	104.2	1
18	73.6	1	99.1	1
19	53.3	1	99.1	0


# Variations on random assignment

- Why is it a good idea to assign the data at random (*i.e.*, shuffle first)?
- Are there any exceptions or alternatives?
  - Our evaluation should match the claims we wish to make about performance
  - We might want to make a specific claim, such as:
    - A model trained on years 2000-2010 generalizes to years 2010-2020
    - A model trained at DUHS generalizes to the UNC Health System
  - We might have repeated measures data
    - Multiple  $(x, y)$  pairs from the same patient *<-- this is a very common source of errors*
    - We typically handle this by assigning repeated measures (*i.e.*, all data for a given patient) to the same fold (*i.e.*, development set, test set)

# Claims about performance must match your evaluation.

## What's wrong here?

- Dataset: pictures of skin lesions along with label (benign/malignant)
  - 2,000 images of 600 lesions from 400 participants
- Goal: predict whether lesions are benign or malignant
- Evaluation Strategy: Divide the data (image/label pairs) at random between a development set (80%) and a test set (20%)



# Claims about performance must match your evaluation.

## What's wrong here?


- Dataset: wearables data from 20 patients with arrhythmia and 20 controls
  - 10 sessions during arrhythmia for each patient with arrhythmia
  - 10 sessions at random times for each control
- Goal: predict arrhythmia
- Evaluation Strategy: Divide the data (image/label pairs) at random between a development set (70%) and a test set (30%)



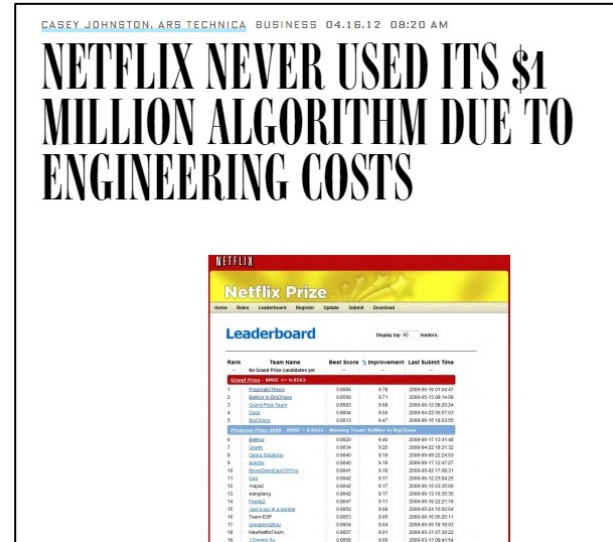
# What goes wrong?

- Subtle mistakes like the previous examples are very common.
- Obvious mistakes like evaluating on the training data are also very common.
- Data Leakage: the use of information in the model training process **which would not be expected to be available at prediction time**
- *Think about whether the information you are using would be available at the time of prediction.*

# Defenses against improper evaluation



Publicly available data and code:  
Facilitates reproducibility



Challenge format:

Test and development sets are fixed by the organizers. Often, test sets aren't even available; entrants submit their model rather than presenting their results.



In healthcare:

There are some good publicly available datasets (MIMIC, *All of Us*, etc.), but we're still way behind the curve.

There's more to model development than training...

What if I want to choose which kind of model is best (*e.g.*, logistic regression, MLP)?

-> this is *model selection*

What if I want to refine or modify my model or the training process in some way (*e.g.*, adjust the number of hidden layers)?

-> this is *hyperparameter tuning*

Definition: a *hyperparameter* is a value or setting that is fixed before training that influences the training process

# Parameters versus Hyperparameters

## Parameters

- Lengths of the seams

## Hyperparameters

- Which seams to adjust
- What ruler to use
- How fast / careless versus careful





# Parameters versus Hyperparameters

## Parameters

- Lengths of the seams
- Model coefficients (*i.e.*, numeric values in the equation)

## Hyperparameters

- Which seams to adjust
- What ruler to use
- How fast / careless versus careful
- Model architecture (*i.e.*, the form of the equation)
- Training algorithm
- Regularization (\*more next time)



# How do we tune hyperparameters?

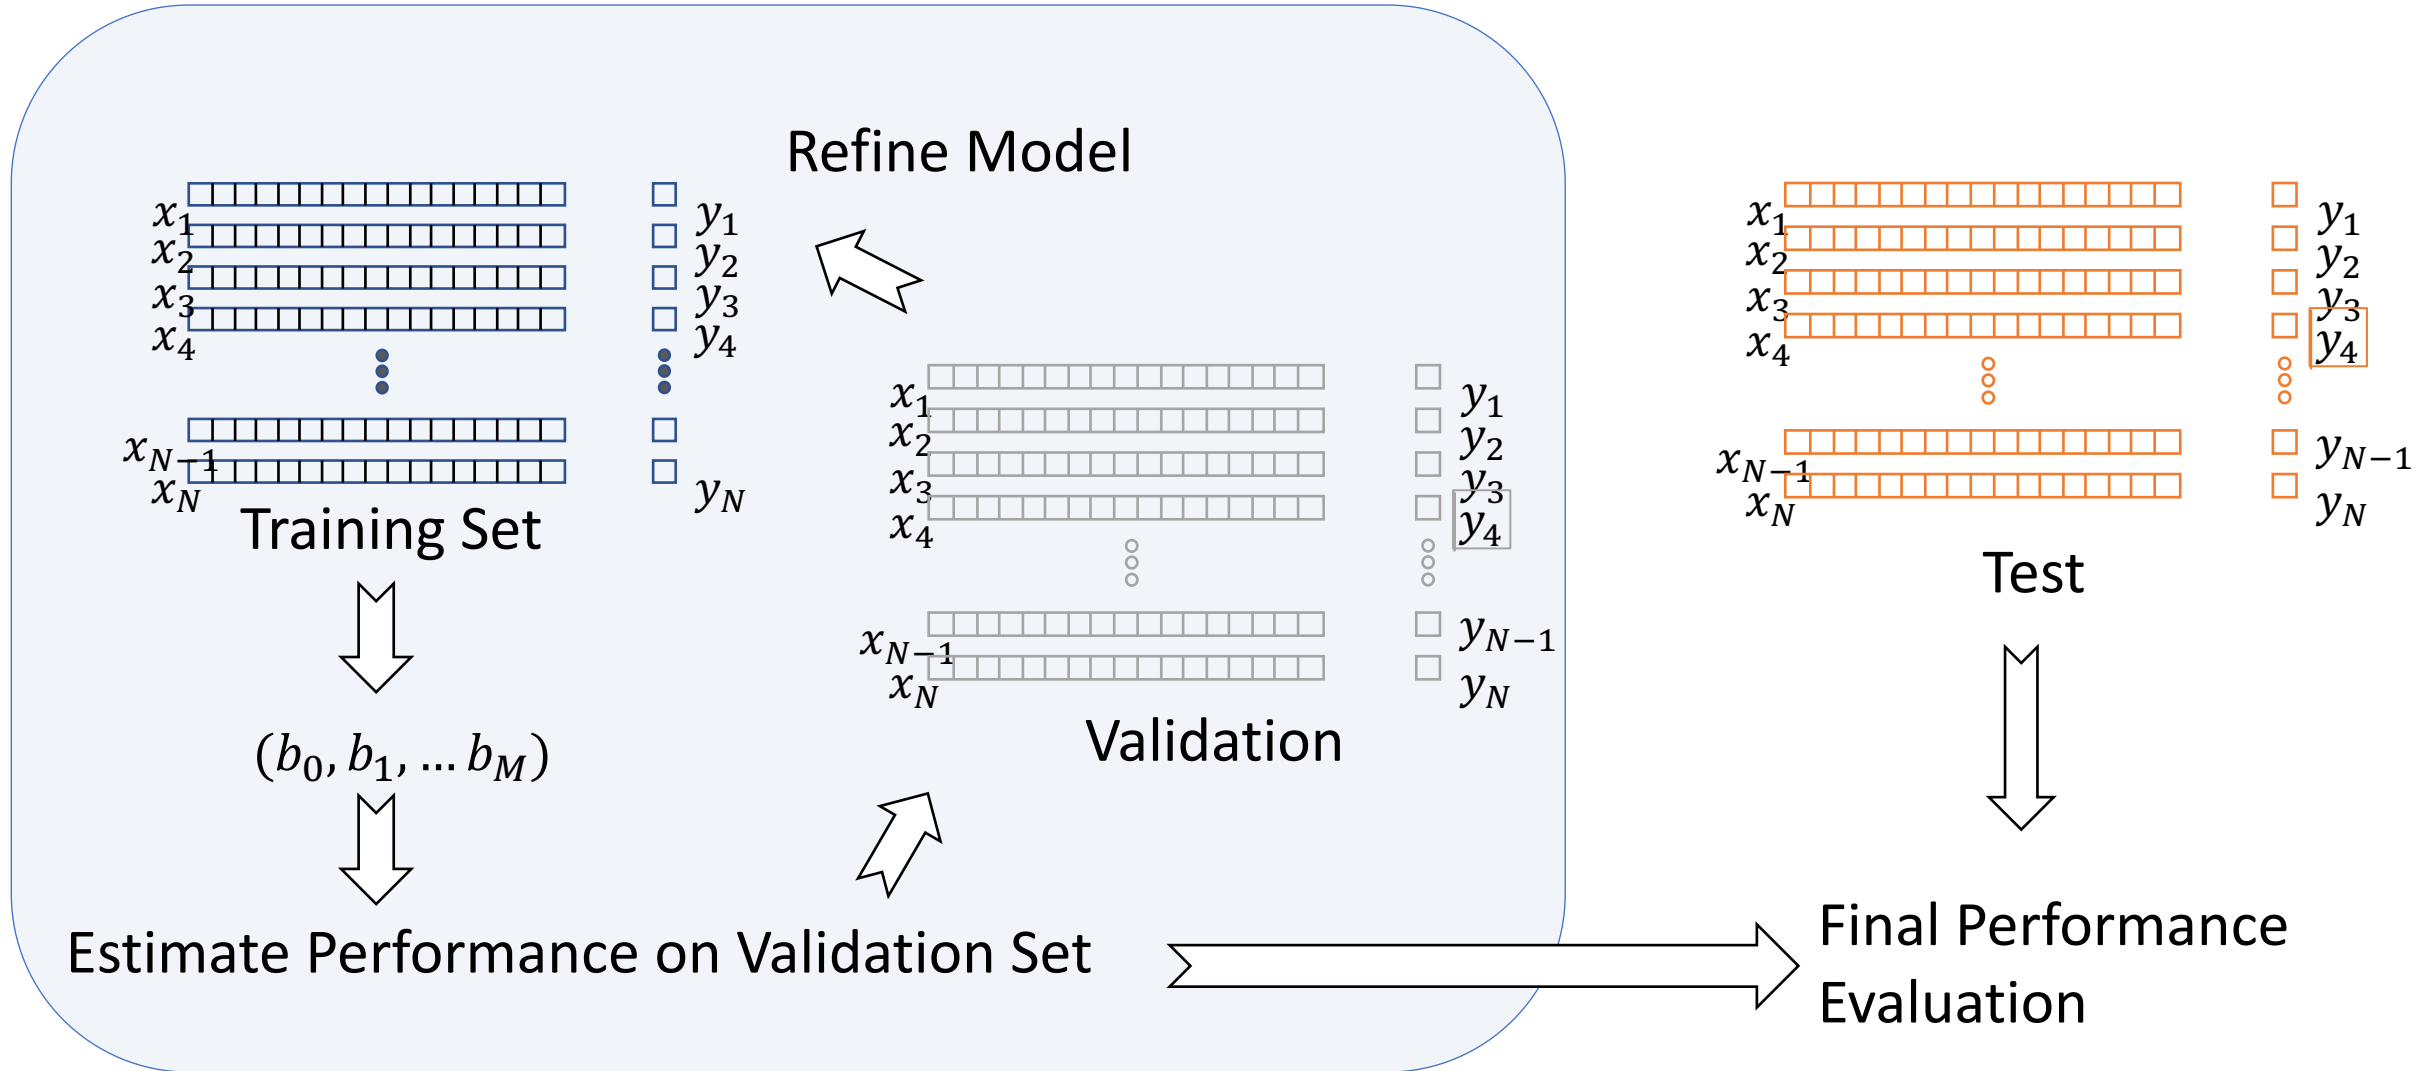
- Proposal 1: Figure out which hyperparameters perform best on the training set
- Proposal 2: Figure out which hyperparameters perform best on the test set
- Are either of these a good idea? Why or why not?

# How do we tune hyperparameters?

- Proposal 1: Figure out which hyperparameters perform best on the training set
- Proposal 2: Figure out which hyperparameters perform best on the test set
- Are either of these a good idea? Why or why not?

-> Hyperparameters require their own *tuning* set.

# How do we divide up the data?




# Cross-Validation

(time permitting)

# Traditional “split” for development + eval



# Cross validation: rotate the test set




Round 1: Train on 1-4, test on 5  
Round 2: Train on all but 4, test on 4  
Round 3: Train on all but 3, test on 3  
Round 4: Train on all but 2, test on 2  
Round 5: Train on 2-5, test on 1

Why would we do this?

- More data for the evaluation
- Better estimate of out-of-sample performance

# What happened to the validation set?

- Use “flat” cross-validation (below) versus “nested” cross-validation
- Both give unbiased estimates and flat is easier



Round 1: Train on 1-3, validate on 4, test on 5

Round 2: Train on 2-4, validate on 5, test on 1

Round 3: Train on 3-5, validate on 1, test on 2

Round 4: Train on 4, 5, and 1; validate on 2, test on 3

Round 5: Train on 5, 1, and 2; validate on 3, test on 4